
When Evaluation Cues Become the Evaluation: Supplementary Material

Anonymous Author(s)

Affiliation

Address

email

1 Evidence and provenance

The primary freeze is the full320-v3 dataset at the pinned revision recorded in provenance/full320-v3/hf_snapshot.json. It contains 11,520 rendered rows from 320 payload blocks. Every activation result used for the main figures is copied into artifacts/pinned_full320/; the submission figure hashes and claim restrictions are recorded in artifacts/submission_figure_manifest.json. The equal-training-size development control is artifacts/full320_equal_n_llama31.8b_l12.json.

2 Complete observational figure set




Figure 1: Held-out purpose AUC across depth for the three pinned model result files.




Figure 2: Cross-format transfer at descriptive observational peaks.




Figure 3: AUC and balanced accuracy at descriptive peaks.

9 3 Equal-N and direct diagnostic controls

10 4 Claim restrictions

11 These artifacts support observational stated-purpose decodability and protocol auditing. They do
 12 not establish evaluation awareness, behavioral steering, causal intervention, confirmation, or a scale
 13 effect. The pilot80-v2 and devbunova datasets are separate freezes and are not pooled with full320-v3.
 14 The equal-N control and direct lens are currently 8B-only.




Figure 4: Four-seed 8B equal-training-size control at layer 12. This is not an all-model replication.




Figure 5: Development-only 8B Jacobian-lens direct diagnostic.