


Cross-format generalization gap is non-monotonic: it collapses near a format-invariant minimum, then re-widens through the late layers

AUC gap (both-format probe - cross-format transfer mean)



Held cue-vocabulary-family OOF · L2 logistic · 7,680 stated rows / 320 payload blocks (3,840 no-cue control rows excluded; no purpose label).

Late-layer re-widening (final-layer gap - minimum-layer gap), 1,000-rep payload-block bootstrap:

Llama-3.1-8B widening 95% CI [0.052, 0.065]; Llama-3.3-70B widening 95% CI [0.063, 0.073]; Llama-3.1-70B widening 95% CI [0.043, 0.052].

Descriptive generalization contrast; not a decomposition of representation content or a causal claim.