

Matched-format projection removes most of the cross-format transfer penalty, while preserving purpose decodability

Post-purpose-formation layers (raw both-format purpose AUC ≥ 0.95), all 3 models.

Llama-3.1-8B Llama-3.1-70B Llama-3.3-70B

Transfer moves far more than purpose decoding does

