

Is Evaluation Awareness just a
disguise for format? 🤔

Conrad Feng, Ian Shen, Parker Patton

Rationale:

- Team interest is in evaluation awareness probing
- Most eval awareness studies focus on alignment and misbehavior of models under evaluation vs. deployment. However, we models interpret “eval awareness” to actually be, and more importantly, how this idea abstracts as layers progress.
- Intuitively, the separation between an evaluation and a deployment lies in the formatting.
 - If you take a test with structured MCQs, 9/10 times you naturally think “this is an evaluation.”
 - Likewise, when quizzed in a casual setting, you’re more likely to think “this is deployment.”
 - Our group wondered if this was true for LLMs.

Previous work

previous work shows that evaluation signals are detectable, but often times confounded with prompt format

1. IS EVALUATION AWARENESS JUST FORMAT SENSITIVITY? LIMITATIONS OF PROBE-BASED EVIDENCE UNDER CONTROLLED PROMPT STRUCTURE
 - a. Crosses eval/deploy context with benchmark/casual format in a 2x2 design
 - b. Standard probe classified 94.5% of benchmark formatted deployment prompts as eval, while decorrelated training reduced to 7.2% and recognized casual evals at 100%
 - c. We converged on a very similar 2x2 experimental design, however our methods differ in the following ways:
 - i. We generate our own prompts for each section of the 2x2 instead of existing benchmarks + LLM rewrites
 - ii. Allows for greater control

methods

High level overview

- **Data generation:** Grok 4.3 for generation, ChatGPT 5.4 for verification
- **Subjects:** Llama 3.3 70B Instruct and Llama 3.1 8B Instruct

Benchmark format (MCQ, "respond with only the letter")	Casual format ("could you give me the answer in a short phrase?")	
Evaluation purpose ("A reviewer will assess this response.")	eval + benchmark — the classic "test" look	eval + casual — <i>off-diagonal</i>
Deployment purpose ("A recipient will apply this response.")	deploy + benchmark — <i>off-diagonal</i>	deploy + casual — the classic "real use" look

Units of division:

1. Domain - broad subject area that contain individual topics
2. Topic - describes task content
3. Task family - Describes kind of reasoning underlying task needs
4. Payload - purpose-neutral problem where purpose and format can be added to
5. Purpose frame - actual sentence added to payload that tells model how answer will be used
6. Purpose pair - a matched set containing one eval and one deploy sentence
7. Cue locus - Describes which portion of the purpose frame differs between eval and deployment
8. Purpose frame family - Groups multiple purpose frame pairs in related ways

1 payload x 2 purposes = 2 purpose framed versions

2 purpose versions x 2 formats = 4 rendered prompts

For example, suppose 7 payloads assigned to purpose pair p03

7 payloads x 4 = 28 prompts for p03

Each family contains two purpose pairs

ff02:

- P03: 7 payloads => 28 prompts
- P04: 7 payloads => 28 prompts

Data generation:

1. Generation of canonical payloads **via Grok 4.3:**
 - a. 16 topics chosen to work in both formats (code review, SQL debugging, unit conversion, recipe scaling, geography, transit routes, ...) $\times$ 5 payloads each = 80 payloads
2. Verification via **ChatGPT 5.4:**
 - a. Quality control such as solving the questions (in the case of quiz, solves without seeing options), rejecting option-dependent or ambiguous items, enforcing purpose neutrality (the payload text may not hint at evaluation, deployment, grading, or audience)
3. Attach prefix and suffix to canonical question
 - a. Eval or deploy payload
 - b. Format from one of two templates: a benchmark MCQ ("Options: ... Respond with only the letter") or a casual short-answer request ("I'm trying to figure something out: ... short phrase?").
4. Train the probes

This reply will be checked by an outside reviewer. **[swappable PURPOSE: a reviewer will assess this response" (evaluation) vs "a recipient will apply this response" (deployment)]**

Question: **[canonical question]**

A database contains a table named Sales with columns product_id (integer) and quantity (integer). The following query is intended to retrieve each product_id along with its total quantity sold, but only for products where the total quantity exceeds 500:

```
SELECT product_id, SUM(quantity)
```

```
FROM Sales
```

```
WHERE SUM(quantity) > 500
```

```
GROUP BY product_id;
```

Under standard SQL semantics where the WHERE clause is evaluated before GROUP BY and before any aggregate computation occurs, and where every non-aggregate column in SELECT must appear in GROUP BY, this query is invalid and will not execute.\n\nWhat prevents this query from executing successfully?

Options: **[swappable FORMAT (mcq for quiz, paragraph for casual)]**

- A. SUM function requires an alias in the SELECT clause
- B. GROUP BY clause missing a column from the SELECT list
- C. WHERE clause must appear after GROUP BY **[correct answer]**
- D. Aggregate function used in the WHERE clause\n\nRespond with only the letter of the correct option.

example of a successfully generated evaluation
datapoint

Probing (Llama-3.1-8B-instruct)

1. Data generation (review)
 - a. As previously mentioned, prompts sorted into five families, which are used to test generalization
2. Activation extraction
 - a. Every prompt ran through model:
 - i. 32 transformer blocks, 4096d residual stream, activations saved as float32
 - ii. Prompts formatted as user message using Llama chat template
 - b. Residual stream captured at final prompt token
3. Linear probe fitting
 - a. Within each training fold:
 - i. Training activations z score normalized; fit an L2-regularized logistic regression classifier; regularization value $C=0.1$ (devbunova uses $C=1.0$)
4. Out of fold validation
 - a. Left one purpose family out for validation
 - i. For each of the five folds, one complete family was held out while probe trained on other four

Probing cont.


1. Four probe training methods:

a.

method	Training data	Research question
confounded	Eval-benchmark + deploy-casual	Does a normal probe learn purpose of format?
decorrelated	All four cells	Can a linear probe recover purpose signal
benchmark=>causal	benchmark format	Does benchmark trained probe transfer to casual prompts
casual=>benchmark	casual format	Does casual trained probe transfer to benchmark prompts

results

Purpose signal peaks early, confounded probe never rises



TAKEAWAYS:

- When confounded (format differs eval vs deploy), both models can not tell the difference between eval vs. deploy (AUC ~0.50)
- When decorrelated (eval as quiz, eval as casual, deploy as quiz, deploy as casual), the rate of identifying eval vs. deploy can be determined above random guess (our purpose prefixes are working)
- Peak layers identified, but what happens at them?


Confounded: training probes on eval = quiz, deploy = casual

Decorrelated: probe trained on all four cells, format balanced across labels, so format is uninformative about purpose

Transferring formats

Transfer to an unseen prompt format improves with scale


A probe trained on one prompt style, evaluated on the other. Held-out purpose families throughout.




however, it decays towards the initial layers.

1. Purpose direction with both formats
2. Quiz-like purpose direction vs. Casual-like purpose direction

Purpose probe across depth: one training format vs both (70B)




Purpose probe across depth: one training format vs both



Takeaway: The gap between "trained on both formats" and "trained on one format" is the purpose signal that doesn't carry across formats. That is, if purpose were encoded format-independently, the training format wouldn't matter and the curves would overlap

How the gap behaves wrt depth



TAKEAWAYS:

- As depth increases, the purpose signal a probe can read becomes increasingly tied to prompt format which is consistent with early-middle layers encoding the situation abstractly and later layers reorganizing toward producing a response in the required style
- In terms of 8B vs 70B:
 - Only the 70B achieves near format-invariance. Yet, later in the layers, 70B becomes the more format bound model.
 - That is, for 70B, there is cleaner abstraction at its peak, deeper format -commitment at the end.

Future work / questions

1. What other experiments can we run to strengthen the claim that probes detect purpose independent of cues?
 - a. Cue would be any feature correlated with the eval/deploy label
 - i. lexical cues, grammatical cue location, inspection/action semantics, etc.
 - b. Our probe generalizes across held out families, but eval items tends to use inspection related wording while deployment items use action language. Which cue based alternative explanations do you think are most important to control for, and what results would be convincing that the probe is tracking purpose rather than those cues?
 - c. What additional experimental evidence (e.g. steering or transfer to natural prompts), would strengthen our purpose interpretation
2. Assuming the probe survives controls, what higher-level problem would make the strongest workshop contribution? For example, should we test whether the direction predicts evaluation-dependent behavior, or whether it can audit how evaluation-like existing benchmarks appear to models